

TM-V71 Remote

User & Technical Manual

Kenwood TM-V71 web remote



Kenwood TM-V71(A/E) web remote + WebRTC audio + HackRF panadapter, for a Raspberry Pi.

Contents

1 Overview	3
2 Features	4
3 Architecture	5
Audio path	5
4 Requirements & Hardware	7
Wiring the USB sound card	7
5 Installation	9
6 Running over HTTPS	10
7 The Web Interface	11
Band panels (VFO A / VFO B)	11
PTT & memory quick keys	11
MUTE and DROP	11
Audio (WebRTC/Opus)	12
HackRF waterfall	13
Selcall (classic 5-tone)	14
Digimodes (CW / RTTY / POCSAG / APRS)	14
APRS	14
Contact columns beside the console	15
Callsign recognition (Vosk)	15
Band scan	16
ASR contacts	17
Voice ID — an over without a callsign	18
Logbook (Wavelog + QRZ.com)	19
Settings	19
8 Mobile App (PWA)	20
9 Trusted Certificate (Root CA)	21
10 Configuration	22
11 REST & WebSocket API	23
Control & status	23
Memory channels	23
Audio	23
Digimodes & selcall	23
SDR & scan	24
Logbook	24
System & power	24
12 Troubleshooting	25
13 Security	26
14 Credits & License	27

1 Overview

TM-V71 Remote is a modern, dependency-light web remote control for the Kenwood TM-V71(A/E) dual-band FM transceiver, built around a direct serial driver. It gives full radio control in the browser, two-way browser audio over WebRTC/Opus, complete memory-channel management, an optional HackRF panadapter, classic 5-tone selective calling, a CW/RTTY/POCSAG digimodes decoder/encoder, a raw RX recorder, and optional off-air callsign recognition (Vosk). It installs as a Progressive Web App (PWA) and is designed to run on a Raspberry Pi.


Unlike hamlib, whose TM-V71 backends are unreliable, this project speaks the radio's documented PC command set directly and exposes the radio's full feature set, including per-channel memory programming.

2 Features

- Full live control of both bands (A/B): frequency, VFO/memory mode, repeater shift & offset, CTCSS/DCS, step, control band, and PTT over CAT.
- Memory channels (CHIRP-level): read, write, delete, rename any of the 1000 channels, plus CSV import/export.
- Live status pushed to the browser over a WebSocket; transmit lights the UI.
- Two-way audio: direct WebRTC/Opus between browser and backend via aiortc; the mic feeds the radio only while PTT is engaged.
- Optional HackRF One waterfall: real-time panadapter (auto-following the tuned frequency) or a wideband sweep.
- Classic 5-tone selective calling (ZVEI-1/2, CCIR, EEA): call, decode, and mute RX until your own ID is received.
- CW (Morse), RTTY (Baudot/AFSK) and POCSAG paging (512/1200/2400 baud, numeric + alphanumeric, BCH FEC) decode + encode over the FM audio path; CW auto mode tracks the received speed and tone pitch, plus a button to decode a captured RX buffer off-line.
- Audio processing: RX de-emphasis (for a flat 9600/discriminator feed), BUSY-gated software squelch, TX AGC, and voice low-pass filters.
- Raw RX recorder with WAV download (e.g. to build ASR training data).
- Off-air callsign recognition (optional, Vosk): detects spoken German callsigns, verifies them against the BNetzA list (name/town/class, or VOID if unassigned), shown in the title bar and a toast.
- Installable PWA with a mobile landscape swipe-deck layout.
- Resilient operation: the phone screen is kept awake, browser audio auto-reconnects after a network glitch, and a backend watchdog releases a latched PTT if every client disappears.
- GPIO power switch, auto power-off, TX power, squelch, in-display S-meter.
- Two themes (dark/light); no build step for the UI.



3 Architecture

Everything below runs on the Pi. The radio is reached over two independent paths — CAT commands on the serial line, and audio on the data port through a USB sound card — and the browser sees only the one HTTPS address the backend serves.



Audio path

One capture callback owns the received audio. What the browser hears is conditioned and squelched; every decoder is fed from a tap taken BEFORE that, on the raw signal, which is why the S-meter, the recorder and both Vosk chains keep working while MUTE or DROP is on. Vosk appears twice on purpose: the grammar-constrained decoder reads spoken callsigns, and a second, plain recognizer with the speaker model turns a whole over into a voiceprint (the grammar decoder runs with N-best alternatives, and in that mode Vosk returns no x-vector).

RX (receive)**Callsign recognition (Vosk, grammar-constrained)****Voice ID (Vosk speaker model)****TX (transmit)**

The decoders always sit on the raw signal — squelch, MUTE and DROP only affect what the browser hears.


The backend owns the serial port directly (backend/app/tmv71.py). One FastAPI process serves the SPA/PWA, the REST control endpoints, the live-status WebSocket, and the WebRTC audio signalling — no extra services. Audio is 48 kHz / 16-bit / mono internally (Opus' native rate).

4 Requirements & Hardware

- Raspberry Pi (tested on Debian 13 / aarch64), Python 3.11+.
- Kenwood TM-V71(A/E) on a serial port (FTDI programming cable), 57600 baud.
- A USB sound interface wired to the radio (data port or mic/speaker), full-duplex.
- System packages: portaudio19-dev, swig + libgpio-dev (optional GPIO).
- Optional: a HackRF One plus the hackrf host tools for the waterfall.
- Optional: vosk + the small German model (offline callsign recognition), and pypdf + the BNetzA Rufzeichenliste PDF (name/town/class + VOID check).

Wiring the USB sound card

The audio does not use one connector but two, and they sit on opposite sides of the radio: transmit audio goes into the mic jack on the side of the control head, received audio comes off the DATA jack on the rear panel. That split is not a preference, it follows from how the radio switches its paths — see the note under the diagram.



TX goes into the mic jack, not into the DATA jack: the radio only routes transmit audio from the DATA jack while a HARDWARE PTT keys it, and this station keys PTT over CAT. RX is taken from PR9, the flat 9600-baud pin: fixed level, untouched by the volume knob, and with neither de-emphasis nor the CTCSS filter in the way — the better feed for APRS. Set menu 519 (PC port baud) to 57600.

- The transmit path needs to lose about 40 dB: the card puts out line level (~1 V), the mic input expects a few millivolts into 600 ohms. A 1:1 transformer in that lead also breaks the ground loop between the Pi's supply and the radio, which is what an alternator whine or a steady hum on your transmitted audio usually is.
- Both receive pins sit at a fixed level, independent of the volume knob, so what the Pi hears does not change when you turn the radio down. They differ in what the radio has already done to the audio. PR9 (pin 4) is the flat output straight off the discriminator; PR1 (pin 5) has been through de-emphasis and the CTCSS filter.
- Take PR9 if APRS matters. De-emphasis tilts the two AFSK tones against each other — 2200 Hz arrives some 5 dB below 1200 Hz — and that twist costs decodes at the edge of the noise; the flat pin hands the demodulator both tones as they were sent. It is also the pin the rest of this program is built around: the software puts the de-emphasis back

for voice (on by default), high-passes the sub-audible CTCSS tone the flat output passes, and gates the audio from the radio's BUSY line instead of relying on the squelch — which on the speaker path would chop a frame in half.

- If PR9 stays silent on your unit, set menu 518 (data speed) to 9600: on some radios that menu switches the receive path as well. It costs nothing here either way, because the transmit audio never touches the DATA jack.
- PTT is keyed over CAT on the PC port, so no PTT line is wired. Nothing in the audio cabling can key the radio.

5 Installation

```
sudo apt-get install -y portaudio19-dev python3-venv swig liblgpio-dev
git clone https://github.com/CQ-DJ0SH/tmv71-remote.git
cd tmv71-remote/backend
python3 -m venv .venv
.venv/bin/pip install -r requirements.txt
```

For a reboot-proof setup, install the systemd unit from the `deploy/` directory. The service runs uvicorn with TLS on port 8443.

6 Running over HTTPS

Browser microphone access (`getUserMedia`) and the PWA service worker require a secure context, so the server runs over HTTPS. A quick self-signed certificate is enough on the desktop (accept the warning once); for installing the PWA on a phone you need a trusted certificate (see chapter 9).

```
cd backend && mkdir -p certs
openssl req -x509 -newkey rsa:2048 -nodes -days 3650 \
  -keyout certs/key.pem -out certs/cert.pem -subj "/CN=tmv71-remote" \
  -addext "subjectAltName=IP:<pi-ip>,DNS:localhost,IP:127.0.0.1"
.venv/bin/uvicorn app.main:app --host 0.0.0.0 --port 8443 \
  --ssl-keyfile certs/key.pem --ssl-certfile certs/cert.pem
```

Open `https://<pi-ip>:8443/` and accept the certificate once.

7 The Web Interface

■ Band panels (VFO A / VFO B)

Each band shows the frequency on a 7-segment display with two stacked meters under a shared S-scale: a real S-meter (S0–S9) in the active band's colour, and below it the AF level / mic-modulation bar (1 s peak-hold). Plus controls for VFO/memory mode, CTRL/PTT band selection, TX power, squelch (remembered per band across power cycles), repeater shift/offset, tone and bandwidth. The digit tuner lets you click bars above/below each digit to step the frequency; AIR Band tunes band A to the 118–137 MHz air band (receive-only).

The S-meter is derived from FM quieting: the high-band noise on the flat RX is inverse to signal strength (loud hiss = no signal, full quieting = strong signal), so it estimates the received signal even though the TM-V71 sends no numeric RSSI over CAT — only a binary BUSY status. It is a relative quieting estimate: a readable signal reads mid/upper scale, a strong local signal saturates at S9, and it snaps back to S0 the moment the carrier drops.

■ PTT & memory quick keys

Hold the large PTT button (or the space bar) to transmit; PTT-LOCK latches transmit. PTT and PTT-LOCK require connected audio (there is no mic otherwise) — they are disabled while audio is off and released automatically if audio disconnects. ROGER adds a descending three-tone beep (d6 · a5 · e5, i.e. 1175/880/659 Hz, 60 ms each with 30 ms gaps) on release, at the level set in the audio settings; while transmitting the button shows a count-up timer (MM:SS). The 1750 Hz button arms a tone-call. Memory quick keys recall channels 0–16 (M0–M9 in the left column, M10–M16 in the right; the loaded channel's key glows); below them the right column sends three DTMF memories (0–2). A status line shows per-band BUSY, the ASR state and the live RX/TX gain (in the PWA; the desktop shows the transmit hint). On mobile, mini RX/TX VU bars with peak-hold flank the button.

■ MUTE and DROP

Two buttons flank the “PTT → BAND x” caption and silence the received audio, differing only in when they stop:

- MUTE (left) is the plain one: it stays until it is pressed again. The label reads MUTED while it is on.
- DROP (right) blanks the transmission in progress and lifts itself as soon as that station stops keying — for an interferer or a rag-chew you do not want to hear, without having to remember to un-mute afterwards.

Both glow in the colour of the transmit band while armed. Three details are worth knowing:

- DROP releases on the falling edge of that band's BUSY signal, not merely because the channel is quiet: arming it on an idle channel keeps the audio muted until the end of the next transmission, rather than un-muting again on the next status update.
- DROP also lifts after three seconds without speech, even while BUSY is still up — the permanent carrier of a repeater would otherwise keep the audio muted indefinitely. The threshold is on the AF level, so a quiet carrier counts as silence while an over does not.

- Muting is done in the browser, on the audio element. The S-meter, the raw RX recorder and the callsign recognition keep receiving the signal — only what you hear is silenced.
- A three-tone chime confirms each change (e5 · a5 · d6 rising on release, the same triad reversed when muting), so the two are told apart without looking. It is generated in the browser and can therefore never be mixed into the transmitted audio, unlike the roger beep, which deliberately is.

■ Audio (WebRTC/Opus)

Open the AUDIO panel, pick the RX band with the RX-A/RX-B switch, click CONNECT and allow the microphone. RX and mic levels are shown live, with the WebRTC RX/TX data rate in the graph corner. Controls: RX/TX gain (with a recommended-default tick), MIC (mic test — meters the mic without keying, records while on and replays your audio over RX when switched off; RX is muted during the test), AGC (automatic TX level), and a small recorder — ● REC / ► PLAY plus a WAV download of the raw, un-squelched RX feed (up to 60 min; e.g. to build ASR training data; the downloaded file is named with the date and time it was saved). Audio timing (TX buffer / TX trail / RX buffer) and the USB card mixer are in Settings > Audio. The link auto-reconnects after a network glitch and is restored on the next launch.

The TX buffer bounds how much microphone audio may queue up for the radio. It has to be larger than the clusters the browser delivers in: measured here with a 40 ms cap, 2.3 s of silence were inserted for want of samples while 2.5 s were discarded at the cap — starving and overflowing at the same time, heard as dropouts. 150 ms is the default and the slider stops at 80 ms, because below that no setting can work. Half of it (40–80 ms) is collected as a cushion before play-out starts.

TX trail is how long the radio stays keyed after the button is released, so the last syllables are not chopped: it covers the frames still in flight over WebRTC at that moment. Once that time is up the over is treated as finished — no further microphone audio is accepted, the silence at the end of it is dropped (below -40 dBFS, with 40 ms kept after the last sound so a soft final consonant survives), and only what was actually spoken plays out before the key goes up. This matters because the browser microphone streams continuously and is not gated by the PTT button: without that cut-off the queue never ran empty and every over was held open for the full timeout, half a second of room noise on the air.

A digital transmission (CW, RTTY, POCSAG, APRS beacon) owns the mic line for as long as it is keyed. Queued microphone audio is discarded when it starts, nothing new is taken while it runs, and the key goes up the moment its last symbol has gone out — with no voice tail and no roger beep, which is a voice-mode habit that would only corrupt the end of the transmission.

The RX buffer (Settings > Audio, 60 ms by default) is a jitter buffer between the sound card and the WebRTC track, and it is worth knowing why it exists. The card produces a 20 ms block on its own clock; the track emits a frame every 20 ms on the backend's clock. Two independent clocks sampling each other means that whenever one runs a few milliseconds early or late, a block goes out twice or is skipped — and every such seam is an audible click. Under load (the callsign recogniser runs continuously, and each over costs another ~1.3 s for its voiceprint) that is frequent enough to sound like crackle. Three 20 ms blocks absorb it; the delay they add is not noticeable in a QSO. Raise it if the received audio still crackles, lower it for the shortest possible delay. Each listening browser gets its own queue,

capped so a client that falls behind loses its oldest blocks instead of accumulating delay; a queue that runs dry refills to the target before playing out again, because limping along one block deep means the next hiccup clicks as well.

RX conditioning (Settings > Audio): a RX de-emphasis (adjustable time constant, on by default) restores natural voice tone when the audio comes from a flat discriminator / 9600-baud data output; a fixed ~180 Hz high-pass on the listen path removes the CTCSS/PL sub-audible tone (67-254 Hz) and DC hum that this flat output passes and de-emphasis would otherwise lift (audible as a low speaker hum), while leaving voice untouched; a BUSY-gated software squelch re-applies muting from the radio's own busy status for that always-open output; and TX/RX voice low-pass filters (≤ 3.5 kHz) tame hiss. The decoders always receive the un-squelched, un-filtered signal.

TX modulation (Settings > Audio). Two stages can be switched into the transmit path, followed by a peak limiter that is always active behind them.

- COMPRESSOR (off by default) raises the AVERAGE modulation: quiet syllables and trailing words come up by as much as 9 dB while loud bursts are held back. Measured on a 1 kHz tone, an input range of $-45\dots-1$ dBFS leaves as $-36\dots-9$ dBFS, so 44 dB of speech dynamics become 27 dB. It acts within a syllable — unlike the AGC, which rides the overall level over seconds; both may be on. Below -52 dBFS nothing is lifted, so the pauses between words keep their room noise in the background instead of being pumped up to speech level.
- PRE-EMPHASIS (on by default) lifts the treble by 6 dB/octave — the exact inverse of the RX de-emphasis, using the same time constant, and within 0.2 dB of the analogue curve (+4.6 dB at 3 kHz). Switch it on ONLY if the transmit audio goes into a FLAT input, i.e. the 9600-baud data port: the radio's mic input and its 1200-baud input pre-emphasise by themselves, and doing it twice sounds shrill. The voice low-pass runs automatically while it is on — without it the lift would send hiss above the voice band as well (+11 dB at 8 kHz, -44 dB with it).
- The LIMITER holds peaks at -1 dBFS whenever either stage is on. It is not decoration: the emphasis lifts treble before the band-limiting low-pass, so a hot microphone clipped in 16-bit — 1.9 % of samples at a -1 dBFS peak — and hard clipping there lands as broadband crackle INSIDE the voice band, where no filter can take it out again. The limiter reduces gain over 2.5 ms slices and ramps between them; measured across the chain, a 0 dBFS input leaves at -1.0 dBFS with not a single clipped sample.

Try either with MIC TEST: the replay runs through the same chain, so it is heard without transmitting.

Bluetooth headsets: transmit audio is captured from the phone's built-in microphone (not the headset's), so the headset stays on the A2DP profile and receive audio keeps coming through in good quality. Using the headset mic would force Android onto the mono HFP/SCO profile and, on many phones, leave RX stuck until Bluetooth is toggled.

HackRF waterfall

If a HackRF One is connected, this panel shows a live spectrum stacked over a waterfall: a panadapter centred on the tuned frequency (auto-following) or a wideband sweep. Receive-only; LNA/VGA gains and a display level are adjustable.

■ Selcall (classic 5-tone)

Send and decode classic selective calls (ZVEI-1/2, CCIR, EEA). Enter a 5-digit CALL code and press CALL (keys PTT). Enter your own ID and press MUTE to silence RX until your ID is received — then it un-mutes automatically. Over FM this is AFSK; use a dummy load when setting up.

■ Digimodes (CW / RTTY / POCSAG / APRS)

Switch between CW (Morse), RTTY (Baudot/AFSK) and POCSAG paging. DECODE shows received text; type into the field and SEND to transmit (keys PTT); the CW text input is forced upper case. Parameters: CW WPM/pitch — with an AUTO mode that tracks both the received speed and tone pitch (shown live on the sliders) and rejects voice/noise so it locks onto the CW even after a spoken ident; RTTY baud/shift/mark; and POCSAG baud (512/1200/2400), RIC, function and numeric/alphanumeric (auto-detected on RX) — e.g. monitor DAPNET on 439.9875 MHz with per-page RIC/FUNC/timestamp output. The REC button decodes the raw RX recorder buffer off-line in the current mode. Over the FM radio this is MCW / AFSK / FSK — not native HF modes.

■ APRS

The fourth digimode, in software: the TM-V71 has no built-in TNC — that is the TM-D710 — so the whole stack runs on the Pi. Feed it the flat 9600-baud data output rather than the speaker path — pin 4 (PR9) of the rear DATA jack, see the wiring diagram in chapter 4: no de-emphasis, and no squelch chopping a frame in half.

- RECEIVE decodes Bell 202 AFSK (1200 baud, mark 1200 Hz, space 2200 Hz) into AX.25 and then APRS: uncompressed and compressed positions, MIC-E — which most mobile stations send — status, messages, objects, telemetry and weather. Each frame becomes one line with source, path, position, course, speed and comment; DEBUG adds a second line with the frame type, symbol, length and the raw payload. The panel head counts frames, damaged frames, distinct stations and noise.
- The CRC check is not optional. Noise regularly produces bit patterns that look like a frame, and every one of them would be a ghost station in the list. Measured here: ten seconds of pure noise produced four such candidates and not one of them survived the check.
- A failed check is two very different things, so they are counted apart. A candidate whose address field still holds twelve plausible callsign characters was a real transmission, damaged on the way — a collision, a fade — and is counted as DAMAGED; with DEBUG on, the line names who was lost. Anything else is noise between two flags and is counted as such. Random bytes pass the address test with a probability around $8 \cdot 10^{-11}$, and sixty seconds of pure noise produced not one. Measured on the air: in 152 s, 17 frames received, 5 damaged, 108 noise candidates — so the raw CRC count said far more about the noise floor than about lost stations.
- TRANSMIT sends a position beacon — BEACON keys PTT, the text field is its optional comment. The position is the centre of the Maidenhead locator from Settings > General, so use six characters or more: four is a $1^\circ \times 2^\circ$ field, roughly 111×70 km. It is sent as an uncompressed position without timestamp, path WIDE1-1, destination APZV71 — the APZ prefix marks experimental software, and claiming a registered one

would misreport which program is on the air. 300 ms of flags precede the frame: too short and the receiving TNC loses the first bytes, which looks like a decoder fault at the far end.

- There is no digipeater, no APRS-IS gateway and no automatic repeat. A beacon that transmits by itself every few minutes is an operating decision, not a default.

Verified by loop-back — every beacon was decoded again by this program's own receiver, position, symbol, path and comment intact, and still error-free at 6 dB signal-to-noise. MIC-E was checked against known coordinates in both hemispheres, including a three-digit longitude with its offset. Decoding costs 0.2 ms per 20 ms of audio.

■ Contact columns beside the console

On a wide screen the console leaves a few hundred pixels unused on either side. They hold a copy of the contact cards — the newest in the left column, the next ones in the right — so the stations heard stay in view while the panel itself is scrolled away. The tray background and the empty places are the panel's own, and the free places are numbered straight through both columns, large and faint, so an empty one says how far the list reaches.

- The heads carry the panel's name and, on the left, the total number of cards — not what fits, so the figure says whether anything is out of sight. The feet carry the two readings worth having in view: the sum of all talk timers on the left, the VOICE MATCH lamp with the number of learned voices on the right.
- The copies are fully operable: delete, log to Wavelog, start and stop the talk timer, mark a card, correct a callsign, and the hover detail. Every click is routed to the original card and performed by ITS button — cloning does not copy event handlers, and giving the copies their own would mean every behaviour existed twice, with the certainty of drifting apart one day.
- One control in the title bar folds both columns away together; they are one display split over two columns, and leaving one open would look like a fault rather than a choice. The columns appear from 1620 px of window width (1220 for the console plus room for both) and never in the mobile deck, which has no margins to put them in.

■ Callsign recognition (Vosk)

An optional, offline speech-recognition pass on the RX audio that detects spoken German callsigns; enable it in Settings > Audio. A grammar-constrained Vosk model (ITU/NATO phonetic alphabet plus German digits — the German spelling alphabet and letter names were dropped as their short, homophone-prone words caused most false matches) stays usable on noisy FM voice; the recognised letters are assembled into a callsign, restricted to the real German BNetzA allocation blocks (always 5–6 characters) and verified against the BNetzA Rufzeichenliste. Accuracy is raised by N-best rescoring — Vosk returns several hypotheses per over and the best callsign across them is chosen, preferring an assigned one — and by repetition voting: a listed call is shown at once, while an unassigned (VOID) hit must be heard twice within a short window before it appears, suppressing one-off mishears (operators send their call 2–3× anyway). A hit appears in a framed field in the title bar (coloured to the RX band) and as a toast, enriched from the offline list with the holder's name, address and licence class (A/E/N); a call that is not assigned is still shown but flagged VOID. Your own callsign is ignored, it runs only while the squelch is open, and it can also

grade the mic-test audio. The callsign list is built once from the register with a converter (python -m app.callsign_list), which splits each entry into callsign, class, name, street, postcode and town, and writes the same list as JSON beside the cache for other tools (--json-only re-exports it in a second, without re-reading the PDF; the TSV stays what the service itself loads, being about 40% faster to read). The register is not consistent about its separator — most entries use a semicolon between holder and address, a few hundred (club and relay stations) only a comma — so a comma in front of the postcode counts as a separator too, and a word hyphenated across a column break is joined again. QRZ.com is used only for a manual lookup in the logbook, never by the ASR.

Region: Germany or the United States

Everything the recognition needs is country-specific, so it is switched in one place: Settings > General > Callsign region. The setting picks the speech model, the words operators spell with, the shape a callsign may have and the register it is checked against. Switching takes a moment — the model is reloaded and the register re-read — and the voice profiles are untouched, the speaker model knowing nothing of language.

- Germany: vosk-model-small-de-0.15, German digits, the BNetzA blocks (DA-DR plus two or three letters), verified against the Rufzeichenliste, classes A / E / N.
- United States: vosk-model-small-en-us-0.15, English digits, a callsign of K, N, W or A with a second letter A-L, then the region digit and one to three letters, verified against the FCC ULS register, classes Novice / Technician / General / Advanced / Extra. A single-letter prefix needs two suffix letters here (K1AB, not K1A): 1x1 calls are special-event and temporary, and allowing four-character reads would hand noise a whole class of false matches.
- Each language has one digit that exists only to survive a noisy channel, and both are in the grammar: "zwo" for 2 in German, "niner" for 9 in English. Leaving them out is not harmless — the word is then absorbed as unknown and the digit drops out of the callsign entirely.

Measured on the bench with synthesised speech, the American chain reads "kilo one juliet tango" as K1JT and pulls Joseph Taylor, Princeton NJ, Extra out of the FCC register; an invented call is flagged VOID as it should be. Two repetitions run together without a pause still merge into one over-long run, exactly as in German — on the air the gap between them separates them.

Powering the radio down suspends the recognition — there is no RX audio to analyse — and switching it back on resumes it. Your setting is not overwritten by this, so detection does not quietly stay off after a power cycle. The ASR indicator and the detected-callsign field stay on screen throughout; only the LED changes: green and pulsing while listening, red while suspended, dim when switched off. The pulse is what says "listening right now", so neither off-state pulses.

Band scan

Sweep a VHF/UHF range or the memory bank and see an occupancy spectrum + waterfall. Double-click a channel to tune the control VFO to it.

■ ASR contacts

A panel below the band scan that collects every recognised station as an index card, so the last overs are readable at a glance instead of as a scrolling log. Cards sit in a tray of empty slots, newest first. The last 200 entries are kept on the Pi and restored when the panel opens; CLEAR ALL empties the view. In the mobile deck the panel has no tab — swipe to it past the band scan.

Each card carries:

- The callsign, in the largest and boldest type on the card — it is the identity of the contact. A zero is drawn slashed (DJØSH) so it cannot be misread as the letter O; that is display only, everything sent to the logbook keeps the plain 0.
- An avatar whose letters and colour are derived from the callsign itself (the characters after the region digit, plus a hue hashed from the whole call), so a station looks the same in every session without anything being stored.
- All three German licence classes A / E / N, with the holder's own one lit and the other two dimmed. If none is lit, the call is not in the BNetzA list.
- Name and town of the holder from the offline BNetzA list — never from QRZ.com, which the ASR does not query. The town carries its postcode in front of it, as on an envelope; the street stays behind the card and travels to the logbook.
- The time the station was last heard and its talk timer, pinned to the bottom edge of the card.

A station heard again does not get a second card. The existing one flashes, is marked, and counts up (×2, ×3 ...) — but it stays where it is, so the tray does not reshuffle under you on every over. Exactly one card carries the red border at a time: the most recently heard. The ordering therefore follows first contact, not last mention.

Hovering a card shows the recogniser detail that used to fill the log lines: word confidence (0.00–1.00, the mean over the individual spelled letters — a statement about the acoustics, not about whether the callsign is real), the S-value and receiving band at the moment of detection, the raw words Vosk actually heard, the rejected N-best candidates, and any 5/6-character correction that was applied.

Two buttons per card: the play symbol logs the QSO straight to Wavelog with the name from the BNetzA list pre-filled, and turns teal once it has gone through so nothing is sent twice. The cross removes a misrecognised contact. That deletion happens on the Pi, not just in the browser: the card is dropped from the log buffer (otherwise it would return the next time the panel opens), every connected client loses it at once, and the callsign is released from the 90-second de-dupe window so a corrected reading can be reported immediately.

Talk time. The clock face in the bottom row of a card is the start/stop button for that station's timer — and the count also runs by itself: it starts on the rising edge of BUSY and stops on the falling one, so an over is timed without touching anything. The clock always belongs to exactly one card, the marked one. Clicking a card moves the mark, and with it the timer, onto that contact — for when the recognition attributed an over to the wrong station. Every card keeps its own total: switching the focus banks the running stretch on the card it came from and picks up the new card's figure, so nothing is lost or counted twice. The display is fixed at hh:mm:ss, otherwise the button would change width as it counts.

The panel head carries the sum of all timers (Σ) and five controls:

- The input field takes a callsign by hand when the recognition misses one — upper-case letters and digits only, filtered while typing. LOG (or Enter) creates the card through the backend exactly like a recognised one, so it lands in the history and reaches every connected client. A message says whether the BNetzA list knew the call.
- MOD marks the focused card as net control. Its timer keeps running and stays readable on the card, but it is left out of every figure: neither the Σ total nor the ranking counts it, and it does not set the scale of the bars — without that, the station holding the round together would flatten every other bar. A flagged card wears a ring around its avatar. The flag lives in the browser (local storage), not on the Pi: it belongs to whoever is listening to this round, and it survives a reload.
- STATS opens the evaluation: every card by talk time, longest first, with callsign, name, a bar and the time. The bar is scaled against the longest over rather than against the sum. Moderators are listed below the ranking, without a bar and without a rank. COPY puts the list on the clipboard as plain text. The dialog keeps counting while a station is being heard, so it can stay open through an over.
- CLEAR ALL empties the log — in dark red, because it takes every card at once, on the Pi and in every open client, and asks before it does. The talk times go with them (they live in the browser); the learned voices do not — clearing a list is not the same statement as “this was wrong”, which is what the cross on a single card means.
- VOICE MATCH, at the right-hand end, switches the speaker recognition on and off (see below). Its lamp pulses green while it is listening; an outlined button means it only observes, a filled one that it acts.

Correcting a callsign: click the call on the card and type over it (Enter applies, Escape and clicking away cancel). The card keeps its talk time, its repeat count and its place, name and town are re-read for the corrected call, and the change reaches every open client. If the corrected call already has a card, the two are merged and their times added. The voice profile is renamed with it — otherwise a mis-heard call would keep collecting a voiceprint under a name that never existed while the real station never gets one.

The interface is in English throughout — labels, hints, messages and error texts. This manual exists in German as well; the labels quoted in it are the English ones, because those are what is on screen.

■ Voice ID — an over without a callsign

Not every over carries a spoken callsign, and the talk timer then counts onto whoever was marked last. Speaker recognition closes that gap: the voiceprint of an over is compared against the stations whose callsign HAS been heard. Enrolment costs nothing — every recognised call labels its own audio, so the profiles build themselves while the panel is simply used.

- An over ends at the PAUSE in speech, not at the carrier. Through a repeater BUSY stays up for the whole QSO, so the falling edge never comes and every station would land in one segment with their voices averaged together. The pause boundary works in simplex too, where the carrier drop simply arrives first.
- Under about 3 s of speech nothing is guessed at: the over is skipped. An over in which two callsigns were heard is skipped as well, because labelling it would spoil both profiles.
- Two stages. Observe only decides and reports — the verdict appears above the card

tray and in the card's hover detail, and nothing moves. Assign additionally moves the mark and the talk timer onto the recognised station; such a card is drawn with a dashed border so it is never confused with the solid red of “just heard saying its call”. The stage is chosen in Settings > Audio > Voice ID.

- A card picked BY HAND always wins: while a manual selection holds, the voice moves neither the mark nor the timer. It is released when the next over begins, so the correction applies to the over it was made in. A callsign actually heard still moves the mark — that is evidence, not a guess.
- It needs the callsign recognition switched on (it listens to the same audio tap), and it costs about 1.3 s of one core per over, afterwards.

Measured on two off-air recordings of the same round (14 + 11 overs, four stations in both): the same station across recordings scored 0.62–0.79 cosine, different stations 0.33 on average and 0.64 at worst. At a threshold of 0.55 with a 0.05 margin over the runner-up, 3 of the 4 known stations were identified, none wrongly, and none of the 7 stations unknown to the profile set was mistaken for a known one. The margin does the real work: a stranger's best match is flat (0.01–0.08 ahead of the second), a true match stands clear (0.20–0.34). Treat it as an assistant that may abstain, not as proof — and never log a QSO on a voice match alone.

The model is Vosk's own speaker model (vosk-model-spk-0.4, ~14 MB), an x-vector trained on 8 kHz telephone speech — narrowband, noisy, channel-varied, which is far closer to FM off a repeater than a wideband model would be. A voiceprint is biometric data: the profiles stay on the Pi in speakers.json (gitignored, never uploaded) and a profile is dropped with its card.

■ Logbook (Wavelog + QRZ.com)

Logs QSOs to a locally installed Wavelog instance. Enter just the callsign (and optionally a name) — frequency, band, mode, date/time and your own callsign are filled in automatically from the live control band and station profile. LOOKUP fetches the operator's name, grid, QTH, country and e-mail from QRZ.com (XML data API) and 'worked before' / DXCC from Wavelog. The ADDRESS field beside QTH takes the holder's street and postcode; logging a contact straight from an ASR card fills it in from the register, and it travels as the ADIF field ADDRESS while QTH keeps the plain town, so Wavelog's statistics stay clean. STATE beside it carries the US state as the ADIF field of the same name — Wavelog counts states for WAS, and a state buried in the address line would not be counted. It stays empty in Germany. LOG QSO sends the contact as ADIF. A green dot shows when Wavelog is reachable; the panel lists the most recent QSOs (with the looked-up details, deletable individually or via CLEAR) and Wavelog's QSO counts (today/month/year/total). Configure the Wavelog URL, API token and station profile, and the QRZ.com username/password, in Settings > Logging. Credentials are stored only on the Pi (runtime.json), never committed.


■ Settings

Tabs: General (callsign, locator, callsign region, API backend URL, serial port/baud, GPIO power, auto power-off, logo, GitHub self-update, Root-CA download), Audio (device, USB mixer, voice filters, test tone, audio timing), Rig-Info, Rig-Memory, Rig-DTMF, Logging (Wavelog + QRZ.com), and Pi-Hardware (host metrics).

8 Mobile App (PWA)

The UI installs as a Progressive Web App: full-screen, with an app-shell service worker for instant launch. On phones the panels become a vertical swipe deck — swipe up/down, one panel per screen (this keeps the deck's scroll axis off the horizontal sliders, so they stay usable) — with the title bar as a slim vertical strip on the left and an icon tab rail on the right. Past the last panel is an info page listing the app version and browser/environment details. The app is forced to landscape; portrait shows a rotate hint. Install via the browser menu (Install / Add to Home Screen); on iOS use Safari > Share.

While the app is open the phone screen is kept awake via the Screen Wake Lock API, so it won't dim or lock mid-QSO. The browser audio link reconnects automatically after a brief network interruption, and if the connection to the backend is lost while transmit is latched, the PTT is released (locally and by a backend watchdog) so the rig can never stay keyed unattended.



9 Trusted Certificate (Root CA)

A self-signed certificate is fine on the desktop, but mobile browsers will not install the PWA or run the service worker without a trusted certificate. Create your own root CA, sign the server certificate with it, and trust the CA on the phone. A download link for the CA appears in Settings > General once it exists.

```
cd backend/certs && mkdir -p ca
# 1) Root CA (10 years) – keep ca.key secret
openssl genrsa -out ca/ca.key 4096
openssl req -x509 -new -key ca/ca.key -sha256 -days 3650 \
  -subj "/CN=TM-V71 Remote Root CA" \
  -addext "basicConstraints=critical,CA:TRUE,pathlen:0" \
  -addext "keyUsage=critical,keyCertSign,cRLSign" -out ca/ca.crt
# 2) server cert signed by the CA (list every name/IP in the SAN)
openssl genrsa -out key.pem 2048
openssl req -new -key key.pem -subj "/CN=tm-v71.example.lan" -out s.csr
openssl x509 -req -in s.csr -CA ca/ca.crt -CAkey ca/ca.key \
  -CAcreateserial -days 825 -sha256 -extfile leaf.ext -out cert.pem
sudo systemctl restart tmv71-remote.service
```

Install ca/ca.crt on the phone (Settings > Security > Install a certificate > CA certificate). The leaf is valid 825 days; re-issue it from the same CA without re-installing on phones. Keep ca/ca.key secret.

10 Configuration

Settings are read from environment variables (prefix TMV71_) or a .env file, with web-UI changes persisted to backend/app/runtime.json. Key variables:

```
TMV71_SERIAL_PORT=/dev/ttyUSB0  TMV71_SERIAL_BAUD=57600
TMV71_HOST=0.0.0.0              TMV71_PORT=8443
TMV71_AUDIO_DEVICE=NAD          TMV71_AUDIO_ENABLED=true
TMV71_GPIO_POWER_PIN=17         TMV71_CALLSIGN=DJ0SH
TMV71_ASR_MODEL_DIR=...         TMV71_ASR_CALLLIST_PDF=...
TMV71_SSL_CERTFILE=...          TMV71_SSL_KEYFILE=...
```

11 REST & WebSocket API

Control & status

GET	/api/status	live radio state
POST	/api/frequency	set VFO frequency
POST	/api/band-mode	VFO / memory / call
POST	/api/control-band	select control band
POST	/api/ptt	key / un-key (CAT)
POST	/api/ptt-band	select TX band
POST	/api/squelch /step	squelch level / step
POST	/api/vfo	shift/offset/tone/bw
GET	/api/info /version	rig + app info
WS	/ws	live status stream

Memory channels

GET	/api/memories?start&end	list channels
GET	/api/memories/{ch}	one channel
PUT	/api/memories/{ch}	write channel
DELETE	/api/memories/{ch}	clear channel
GET	/api/memories.csv	export CSV
POST	/api/memories/import	import CSV
POST	/api/recall	recall to band

Audio

POST	/api/webrtc/offer	WebRTC SDP offer
GET	/api/audio/status	RX/TX levels, flags
GET	/api/audio/devices	list sound devices
POST	/api/audio/device	pick device
POST	/api/audio/gain	rx/tx gain + TX AGC
POST	/api/audio/buffer	tx buffer / ptt tail / rx buffer
POST	/api/audio/tones	roger/test/mic/lowpass/de-emph/squelch
POST	/api/audio/record[/clear]	raw RX recorder
GET	/api/audio/record.wav	download recording (WAV)
GET/POST	/api/audio/mixer	USB card mixer

Digimodes & selcall

GET	/api/digi	CW/RTTY/POCSAG/APRS status
POST	/api/digi/config	mode + parameters
POST	/api/digi/tx	encode + transmit
POST	/api/digi/decode-recording	decode the RX buffer
WS	/ws/digi	decoded text stream
GET/POST	/api/asr/config	callsign recognition (Vosk)
POST	/api/asr/log	add a callsign by hand
PATCH	/api/asr/log/{call}	correct a callsign (profile follows)
GET/POST	/api/asr/speaker	voice ID: on/off, stage, limits
DELETE	/api/asr/log	empty the contact log (CLEAR ALL)
DELETE	/api/asr/log/{call}	drop a misrecognised contact

WS	/ws/callsign	recognised-callsign events
GET	/api/selcall	5-tone status
POST	/api/selcall/config	standard / tone / own
POST	/api/selcall/tx	send a 5-tone call
WS	/ws/selcall	decoded calls stream

SDR & scan

GET	/api/hackrf	SDR status
POST	/api/hackrf/start stop config	
WS	/ws/hackrf	spectrum/waterfall frames
GET	/api/scan	POST /api/scan/start stop band scan

Logbook

GET/POST	/api/log/config	logbook credentials (Wavelog + QRZ)
POST	/api/log/test	test the Wavelog connection
POST	/api/log/qrz/test	test the QRZ.com login
GET	/api/log/stations	Wavelog station profiles
POST	/api/log/lookup	callsign lookup (QRZ + Wavelog)
POST	/api/log/qso	log a QSO
GET	/api/log/recent	recent QSOs + online + Wavelog stats
POST	/api/log/recent/delete	delete one recent entry
POST	/api/log/recent/clear	clear the recent list

System & power

GET/POST	/api/power-switch	GPIO power
POST	/api/gpio-config	set GPIO pin
POST	/api/auto-power-off	idle auto-off
GET/POST	/api/serial-config	serial port/ baud
GET/POST	/api/callsign /theme	
GET	/api/system	Pi host metrics
GET/POST	/api/update	GitHub self-update

12 Troubleshooting

- No CAT / 'radio offline': check the serial port and baud in Settings; the FTDI cable must be on /dev/ttyUSB0 (or set the right port).
- No audio: ensure HTTPS, click CONNECT and allow the mic; check the USB card and its mixer levels (playback drives the radio mic on TX).
- PWA won't install / theme not switching on a phone: the certificate is not trusted — install the Root CA (chapter 9).
- RX filter only on one channel: reconnect audio (Disconnect/Connect) to renegotiate Opus to mono.
- Decoders need a clean signal; tune levels and (for RTTY) the mark tone.

13 Security

LAN-only by design; there is no authentication. Do not expose the port directly to the internet — use a VPN (WireGuard/Tailscale) or a reverse proxy with TLS + auth. The Root CA private key never leaves the Pi.

14 Credits & License

Kenwood PC protocol docs: LA3QMA/TM-V71_TM-D710-Kenwood. Built with aiortc (WebRTC/Opus) and sounddevice. Fonts: Saira, IBM Plex Mono, DSEG (7-segment), Neuropol (title). See the repository for license details.